

Homework 3

Solutions

2026-01-08

Introduction

In this assignment you will practice (1) choosing appropriate plots based on variable type and analytic goal, and (2) using `dplyr` pipelines to move from raw data → transformed data → a final plot.

You will use two data sets:

- `ncbirths` (from the `openintro` package)
- `flights` (from the `nycflights13` package)


Use the following code chunk to load the necessary packages (`tidyverse`, `ggpubr` and `sjPlot`), and load the data listed above.

```
library(tidyverse)
library(ggpubr)
library(sjPlot)
library(gtsummary)
flights <- nycflights13::flights
ncbirths <- openintro::ncbirths
```

Part I: Creating plots (ncbirths)


1. Create appropriate visualizations for the `marital` status, and the mothers age (`mage`).
 - a. `marital` status

```
ggplot(ncbirths, aes(x = marital)) + geom_bar()
plot_frequencies(ncbirths$marital)
table(ncbirths$marital) |> pie()
```



b. Mothers age (mage)


```
ggplot(ncbirths, aes(x = mage)) + geom_histogram()
ggplot(ncbirths, aes(x = mage)) + geom_density()
ggplot(ncbirths, aes(x = mage)) + geom_boxplot()
gghistogram(ncbirths, x="mage")
ggdensity(ncbirths, x="mage")
ggboxplot(ncbirths, y="mage")
```



2. Bivariate Relationships


- a. Is birth weight related to length of pregnancy? Create a scatterplot of `weight` versus `weeks`. Add **one** smoother to your plot (either a loess smoother *or* a linear model line).

```
ggplot(ncbirths, aes(x = weeks, y = weight)) +
  geom_point() + geom_smooth(se = FALSE)
ggscatter(ncbirths, x="weeks", y="weight", add = "loess")
```



- b. Do babies of smokers tend to have different birth weights than babies of non-smokers?
Create a plot that compares the distribution of weight across levels of habit.

```
ggplot(ncbirths, aes(x = weight, fill = habit)) + geom_histogram()
ggplot(ncbirths, aes(x = weight, fill = habit)) + geom_density(alpha = .3)
ggplot(ncbirths, aes(x = habit, y = weight, fill = habit)) +
  geom_boxplot()
gghistogram(ncbirths, x="weight", fill = "habit")
ggdensity(ncbirths, x="weight", fill = "habit")
ggboxplot(ncbirths, x = "habit", y = "weight", fill = "habit")
```



- c. Does smoking habit differ by mothers maturity status? Create a table to display the distribution of smoking habit by the mothers maturity status (`mature`).

Characteristic	nonsmoker N = 873 ¹	smoker N = 126 ¹
mature		
mature mom	121 (92%)	11 (8.3%)
younger mom	752 (87%)	115 (13%)

¹n (%)


```
table(ncbirths$mature, ncbirths$habit) |> prop.table(margin=1) |> round(3)
```

	nonsmoker	smoker
mature mom	0.917	0.083
younger mom	0.867	0.133

```
tbl_summary(ncbirths, include = "mature", by = "habit", percent = "row")
```

d. Create a side by side bar chart that visually reflects the comparison above.

```
ggplot(ncbirths, aes(x=mature, fill=habit)) + geom_bar(position = "dodge")
plot_xtab(x = ncbirths$mature, grp = ncbirths$habit, show.total = "false",
margin = "row")
```



Part II: Efficient Data Management → Plot (NYC Flights)

Each question starts with the full `flights` data set and will use multiple `dplyr` verbs so chaining is advised. Your final result should be shown as a printed tibble.

1. Create a new data set that contains only flights that:

- departed from **JFK**
- traveled at least **1000 miles**
- have non-missing `air_time`

Keep only the variables `origin`, `dest`, `distance`, and `air_time`. Then create a new variable called `speed` defined as `distance / air_time * 60`. Finally, sort from fastest to slowest speed and display the first 10 rows.

```
flights %>%
  filter(origin == "JFK", distance >= 1000, !is.na(air_time)) %>%
  select(origin, dest, distance, air_time) %>%
  mutate(speed = distance / air_time * 60) %>%
  arrange(desc(speed)) %>%
  head(10)
```

```
# A tibble: 10 x 5
  origin dest  distance air_time speed
  <chr>  <chr>    <dbl>     <dbl>   <dbl>
1 JFK    SJU      1598      170    564
2 JFK    SJU      1598      172    557.
3 JFK    STT      1623      175    556.
4 JFK    SJU      1598      173    554.
5 JFK    SJU      1598      173    554.
6 JFK    SJU      1598      173    554.
7 JFK    SJU      1598      173    554.
8 JFK    SJU      1598      173    554.
9 JFK    SJU      1598      173    554.
10 JFK   SJU      1598      173    554.
```

2. What are the most popular destinations for each airport?

This question is broken down into a few smaller parts to help you ensure that you are on the right track before moving to the next step. Only add one command at a time and “trust but verify” to make sure your code works as intended. You do not need to report the results of each step, just the final result.

- a. Use `filter()` to keep only flights with a non-missing distance, then `group_by(origin, dest)` to group the data by origin first, then destination within origin.
- b. Use `summarise()` to calculate, for each origin–destination pair:
 - the number of flights as `n_flights`

- the average distance as `avg_distance`
- Keep only the top 3 destination using `slice_head(n=3)` after sorting by the number of flights.
 - Use `relocate()` so the columns appear in the order: `origin, avg_distance, dest, n_flights`, and display the resulting table.

```
flights %>%
  filter(!is.na(distance)) %>%
  group_by(origin, dest) %>%
  summarise(
    n_flights = n(),
    avg_distance = mean(distance)) %>%
  arrange(desc(n_flights)) %>%
  slice_head(n=3) %>%
  relocate(c(dest, n_flights), .after = last_col())
```

```
# A tibble: 9 x 4
# Groups:   origin [3]
  origin avg_distance dest   n_flights
  <chr>      <dbl> <chr>     <int>
1 EWR          719  ORD       6100
2 EWR          200  BOS       5327
3 EWR          2565 SFO       5127
4 JFK          2475 LAX      11262
5 JFK          2586 SFO       8204
6 JFK          187  BOS       5898
7 LGA          762  ATL      10263
8 LGA          733  ORD       8857
9 LGA          544  CLT       6168
```

3. Exploring travel delays

Typically we don't care if a departure delay is a few minutes late, but anything longer could be detrimental to connecting flights. Use `case_when` to bin the `dep_delay` variable into categories based on departure delay status.

- Create **at least three categories**
- Choose and briefly describe your cutoff values
- Use category names that correspond to intuitive delay types (e.g., “on time”, “short,” “medium,” “long”).
- A reasonable approach is to look at the distribution of the departure delay first (histogram), then choose cutoffs based on where the distribution naturally changes.

Characteristic	N = 336,776 ¹
delay_category	
long delay	48,291 (15%)
on time or early	200,089 (61%)
short delay	80,141 (24%)
Unknown	8,255
¹ n (%)	

Display a frequency table showing how many flights fall into each delay category.

```
flights %>%
  mutate(delay_category = case_when(
    dep_delay <= 0 ~ "on time or early",
    dep_delay > 0 & dep_delay <= 30 ~ "short delay",
    dep_delay > 30 ~ "long delay"
  )) %>%
 tbl_summary(include = delay_category)
```