

Working with Data Frames - Notes

Key

2025-12-26

! How to use the Notes files

This notes file follows along with the lesson materials. Your job is to fill out the missing code, often denoted with _____, and to complete the “you try it” type problems presented in the notes.

Render this document regularly, ensuring that it is complete before you submit. Before you submit, remove the `error: true` line in the YAML header then render.

Setup

```
library(ggplot2)          # replace the blank line with `ggplot2`  
library(gtsummary)  
ncbirths <- openintro::ncbirths
```

Frequency Tables

You can create a basic frequency table by using the `table()` function.

```
table(ncbirths$lowbirthweight)
```

```
low not low  
111     889
```

Relative frequencies (proportions or percentages) are calculated by putting the results of the `table` function inside the `prop_table` function.

Characteristic	N = 1,000 ¹
visits	12 (10, 15)
Unknown	9
lowbirthweight	
low	111 (11%)
not low	889 (89%)

¹Median (Q1, Q3); n (%)

```
prop.table(
  table(ncbirths$lowbirthweight)
)
```

```
low not low
0.111 0.889
```

Summary Statistics

The function `summary()` prints out the five number summary, and includes the mean. This function also displays the number of missing values for that variable.

```
summary(ncbirths$visits)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
0.0	10.0	12.0	12.1	15.0	30.0	9

Fancy summary tables

The `gtsummary` package provides a single function `tbl_summary` to create a really nicely formatted summary table for both quantitative and categorical data types.

```
tbl_summary(ncbirths,
  include = c(visits, lowbirthweight)
)
```

The `statistic` argument can be used to change what values are displayed.

Characteristic	N = 1,000 ¹
visits	12 (4)
Unknown	9
lowbirthweight	
low	111 / 1,000 (11%)
not low	889 / 1,000 (89%)

¹Mean (SD); n / N (%)

```
tbl_summary(ncbirths,
            include = c(visits, lowbirthweight),
            statistic = list(
                all_continuous() ~ "{mean} ({sd})",
                all_categorical() ~ "{n} / {N} ({p}%)"
            )
        )
```

Missing Data

We can see 4 out of the first 6 values for the variable `fage` (fathers age) in the `ncbirths` data set are missing.

```
head(ncbirths$fage)
```

[1] NA NA 19 21 NA NA

Problem 1: R can't do arithmetic on missing data.

```
mean(ncbirths$fage)
```

[1] NA

Fix this error

Add the argument `na.rm=TRUE` to the `mean()` function inside the right hand parenthesis to calculate the mean after excluding missing values.


```
mean(ncbirths$fage, na.rm = TRUE)
```

```
[1] 30.25573
```

Problem 2: Some plots will show NA as it's own category

Sometimes this is fine, other times this is undesirable. We'll see later how we can adjust this plot to remove that column of NA.

```
ggplot(ncbirths, aes(premie)) + geom_bar()
```



Identifying missing values

To find out how many values in a particular variable are missing we can use several different approaches.

Look at the raw data

```
head(ncbirths)
```

```
# A tibble: 6 x 13
  fage   mage mature    weeks premie visits marital gained weight lowbirthweight
  <int> <int> <fct>     <int> <fct>   <int> <fct>   <int> <dbl> <fct>
1    NA     13 younger ~     39 full ~     10 not ma~     38    7.63 not low
2    NA     14 younger ~     42 full ~     15 not ma~     20    7.88 not low
3    19     15 younger ~     37 full ~     11 not ma~     38    6.63 not low
4    21     15 younger ~     41 full ~      6 not ma~     34     8    not low
5    NA     15 younger ~     39 full ~      9 not ma~     27    6.38 not low
6    NA     15 younger ~     38 full ~     19 not ma~     22    5.38 low
# i 3 more variables: gender <fct>, habit <fct>, whitemom <fct>
```

Look at data summaries

Functions such as `table()` have a `useNA="always"` option to show how many records have missing values, and `summary()` will always show a column for NA.

```
table(ncbirths$habit, useNA="always")
```

nonsmoker	smoker	<NA>
873	126	1

```
summary(ncbirths$fage)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
14.00	25.00	30.00	30.26	35.00	55.00	171

Use a logical statement

The function `is.na()` returns TRUE or FALSE for each element in the provided vector for whether or not that element is missing.

```
x <- c("green", NA, 3)
is.na(x)
```

```
[1] FALSE  TRUE FALSE
```

```
sum(is.na(ncbirths$fage))
```

```
[1] 171
```

There are 171 records in this data set where the age for the father is not present.

i Negating `is.na()` to find the non-missing values

Sometimes you want to operate only on the non-missing values. Recall from [?@se-logical](#) we can use the `!` to negate a boolean argument.


```
!is.na(x)
```

```
[1] TRUE FALSE TRUE
```

Hide the NA bar

We can use this tactic to fix that barchart from above. Wrap `!is.na()` around the `premie` variable in the `ggplot` code to create a barchart that does not have a bar for missing values.

```
ggplot(ncbirths, aes(!is.na(premie))) + geom_bar()
```



Data management

Overwrite existing values

💡 Example 1: Too low birthweight

Let's look at the numerical distribution of birthweight (in pounds) of the baby.

```
summary(ncbirths$weight)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.000	6.380	7.310	7.101	8.060	11.750

The value of 1 lb seems very low. The researchers you are working with decide that is a mistake and should be excluded from the data. We would then set all records where `weight=1` to missing.

```
ncbirths$weight[ncbirths$weight==1] <- NA
```

```
min(ncbirths$weight, na.rm=TRUE)
```

```
[1] 1.19
```

Your Turn

But what about other weights that aren't quite as low as 1, but still unusually low?

- Write the code to set all birth weights less than 4 lbs (<4) to missing (NA).
- Then recalculate the mean to confirm your recode worked.

```
ncbirths$weight[ncbirths$weight < 4] <- NA  
min(ncbirths$weight, na.rm=TRUE)
```

```
[1] 4
```

Creating new variables

The following code creates a new variable `wtgain_mom` the weight gained by the mother, that is not due to the baby by subtracting `weight` from `gained`.

```
ncbirths$wtgain_mom <- ncbirths$gained - ncbirths$weight
```

To confirm this variable was created correctly, we look at the data contained in three variables in question.

```
head(ncbirths[,c('gained', 'weight', 'wtgain_mom')])
```

```
# A tibble: 6 x 3
  gained weight wtgain_mom
  <int>   <dbl>      <dbl>
1     38    7.63     30.4
2     20    7.88     12.1
3     38    6.63     31.4
4     34     8        26
5     27    6.38     20.6
6     22    5.38     16.6
```

Dichotomizing data

Let's add a variable to identify if a mother in the North Carolina births data set was underage at the time of birth. Specifically Make a new variable `underage` on the `ncbirths` data set. If `mage` is under 18, then the value of this new variable is `underage`, else it is labeled as `adult`.

```
ncbirths$underage <- ifelse(ncbirths$mage < 18, "underage", "adult")
```

Trust but Verify

```
table(ncbirths$underage, useNA="always")
```

adult	underage	<NA>
963	37	0

Next let's check it against the value of `mage` itself. Let's look at all rows where mothers age is either 17 or 18 `mage %in% c(17,18)`, and only the columns of interest.

```
ncbirths[ncbirths$mage %in% c(17,18),c('mage', 'underage')]
```

```
# A tibble: 57 x 2
  mage underage
  <int> <chr>
1    17 underage
2    17 underage
3    17 underage
4    17 underage
5    17 underage
6    17 underage
7    17 underage
8    17 underage
9    17 underage
10   17 underage
# i 47 more rows
```

Chaining commands

Example: Frequency tables & summary statistics using the pipe

First stating the variable, then pipe in the summary function.

```
ncbirths$mature |> table() # instead of table(ncbirths$mature)
```

```
mature mom younger mom
      133          867
```

```
ncbirths$mage |> mean() #instead of mean(ncbirths$mage)
```

```
[1] 27
```